D1 Overall system and trust boundaries

The Python backend authorizes and executes side effects; Rust is the desktop shell only.

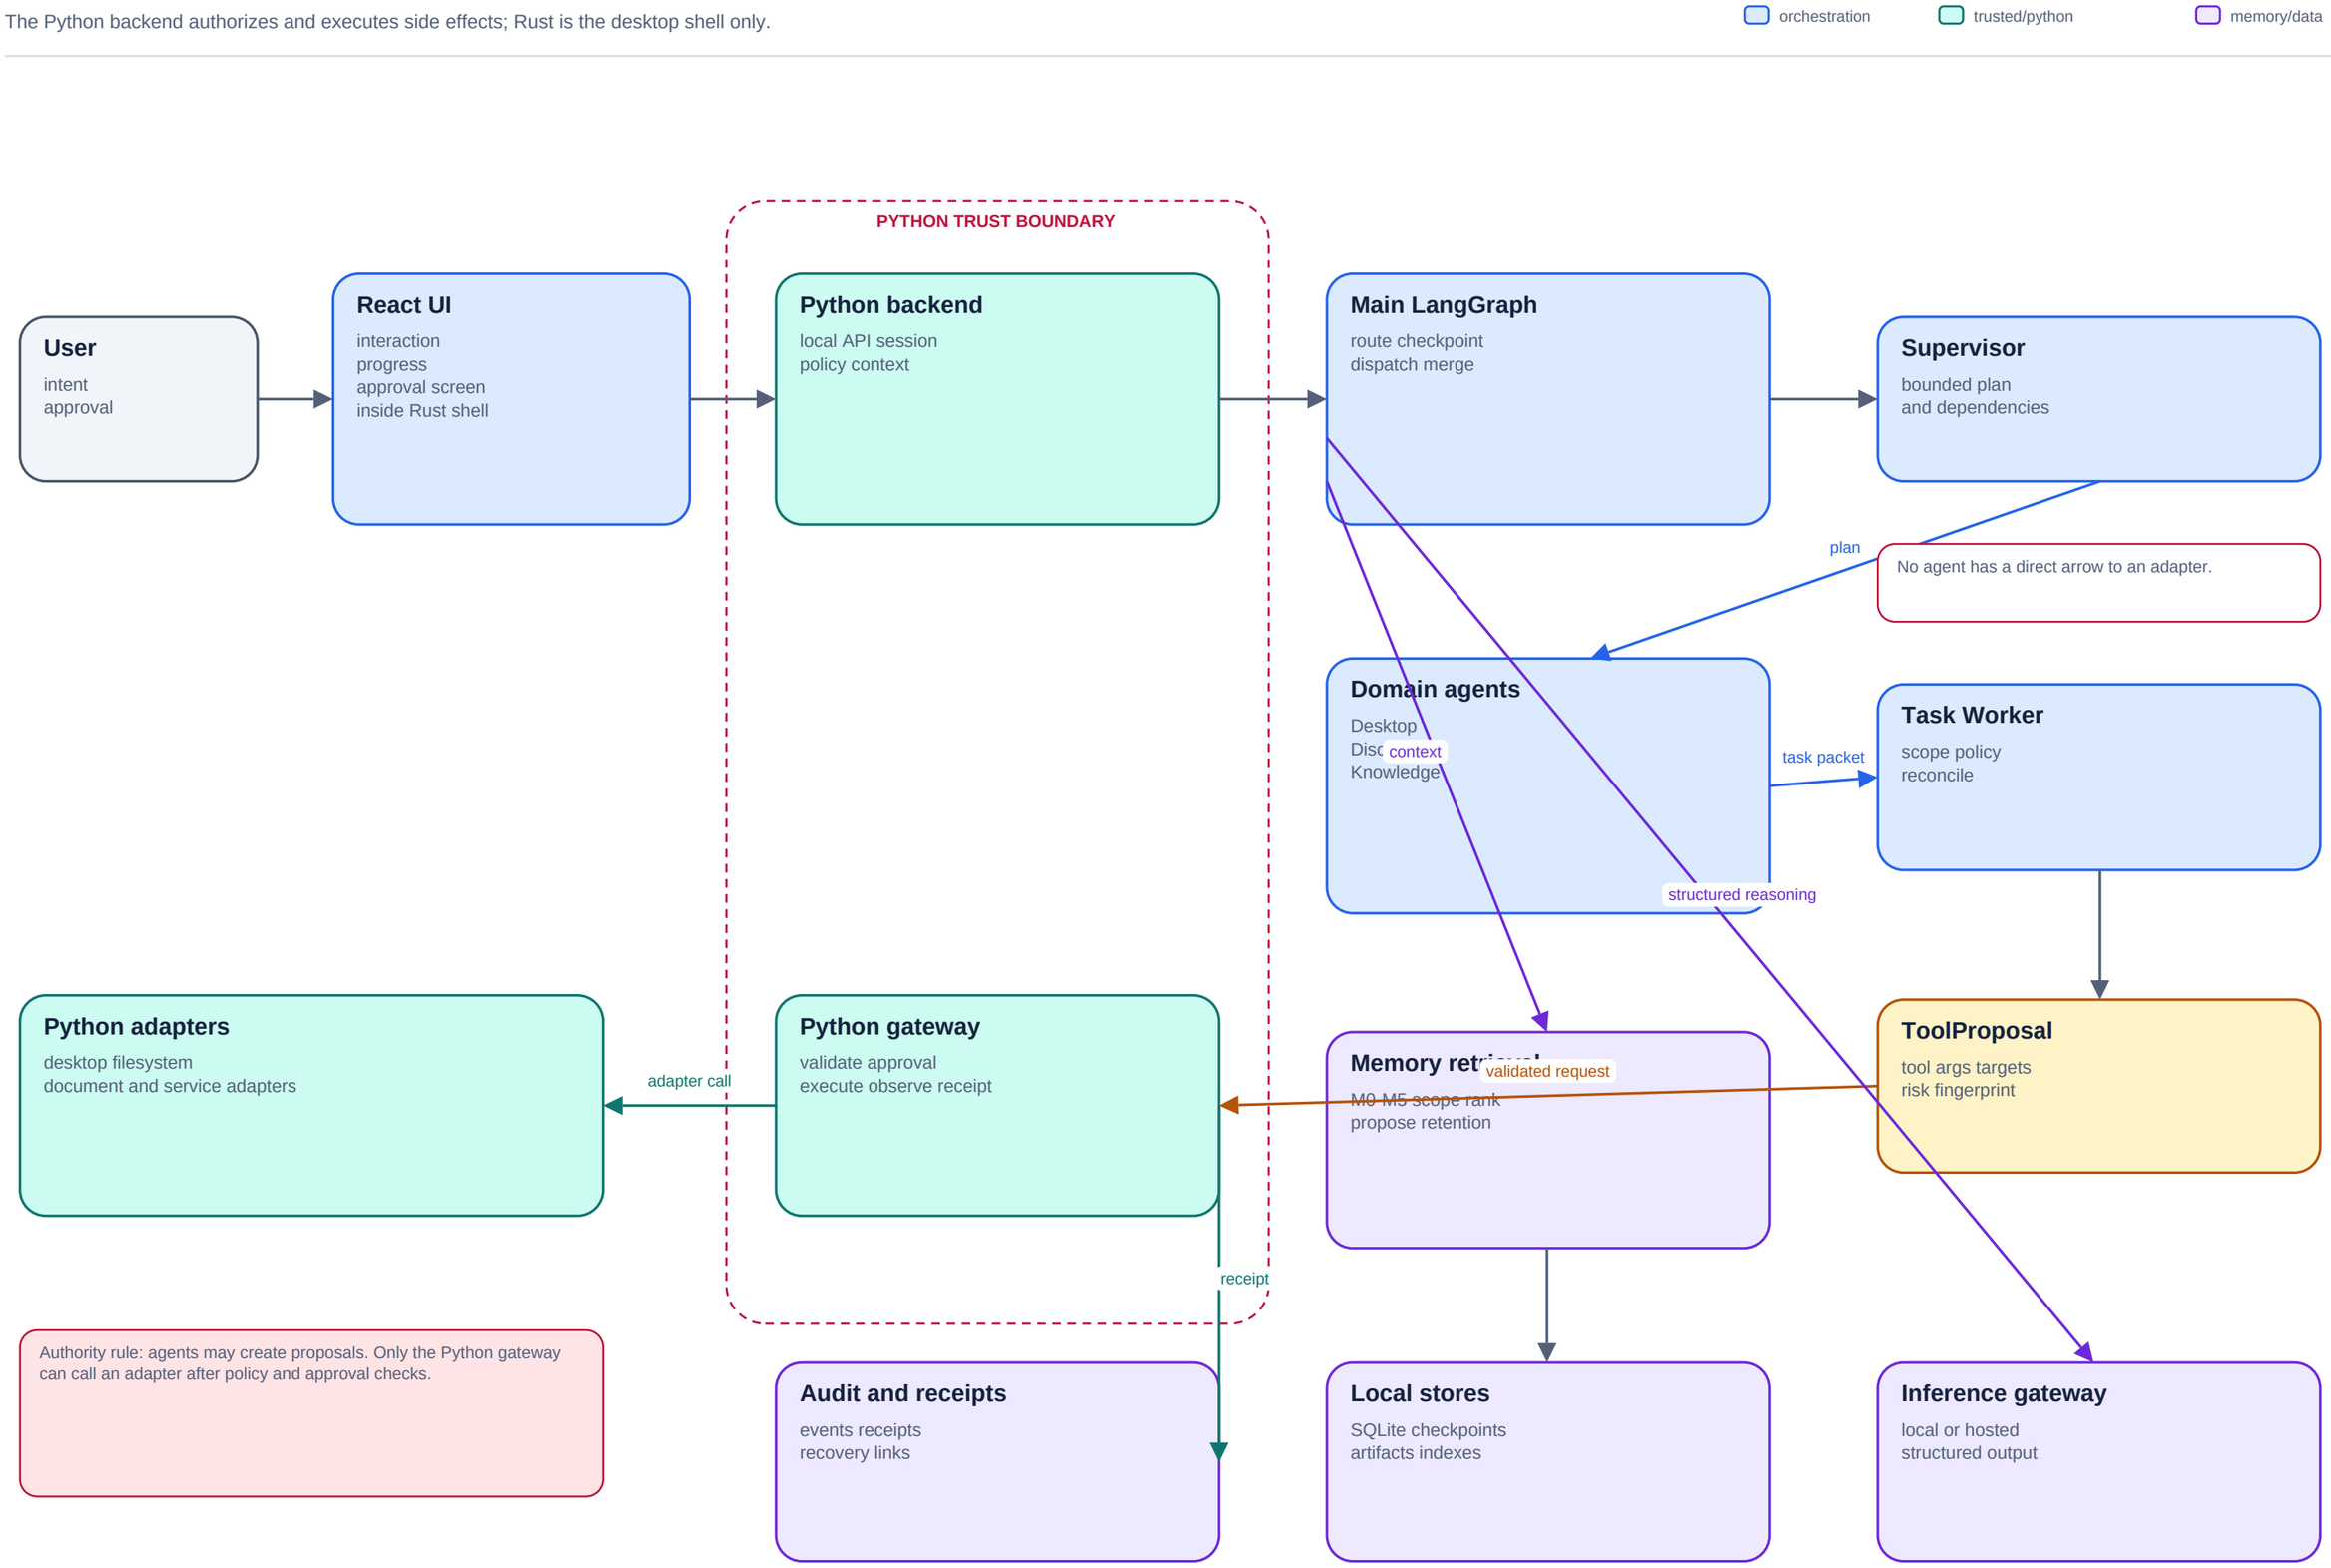

D2 Main LangGraph execution flow

Exact source-supplied node names; arrows show a bounded, recoverable proposed routing model.

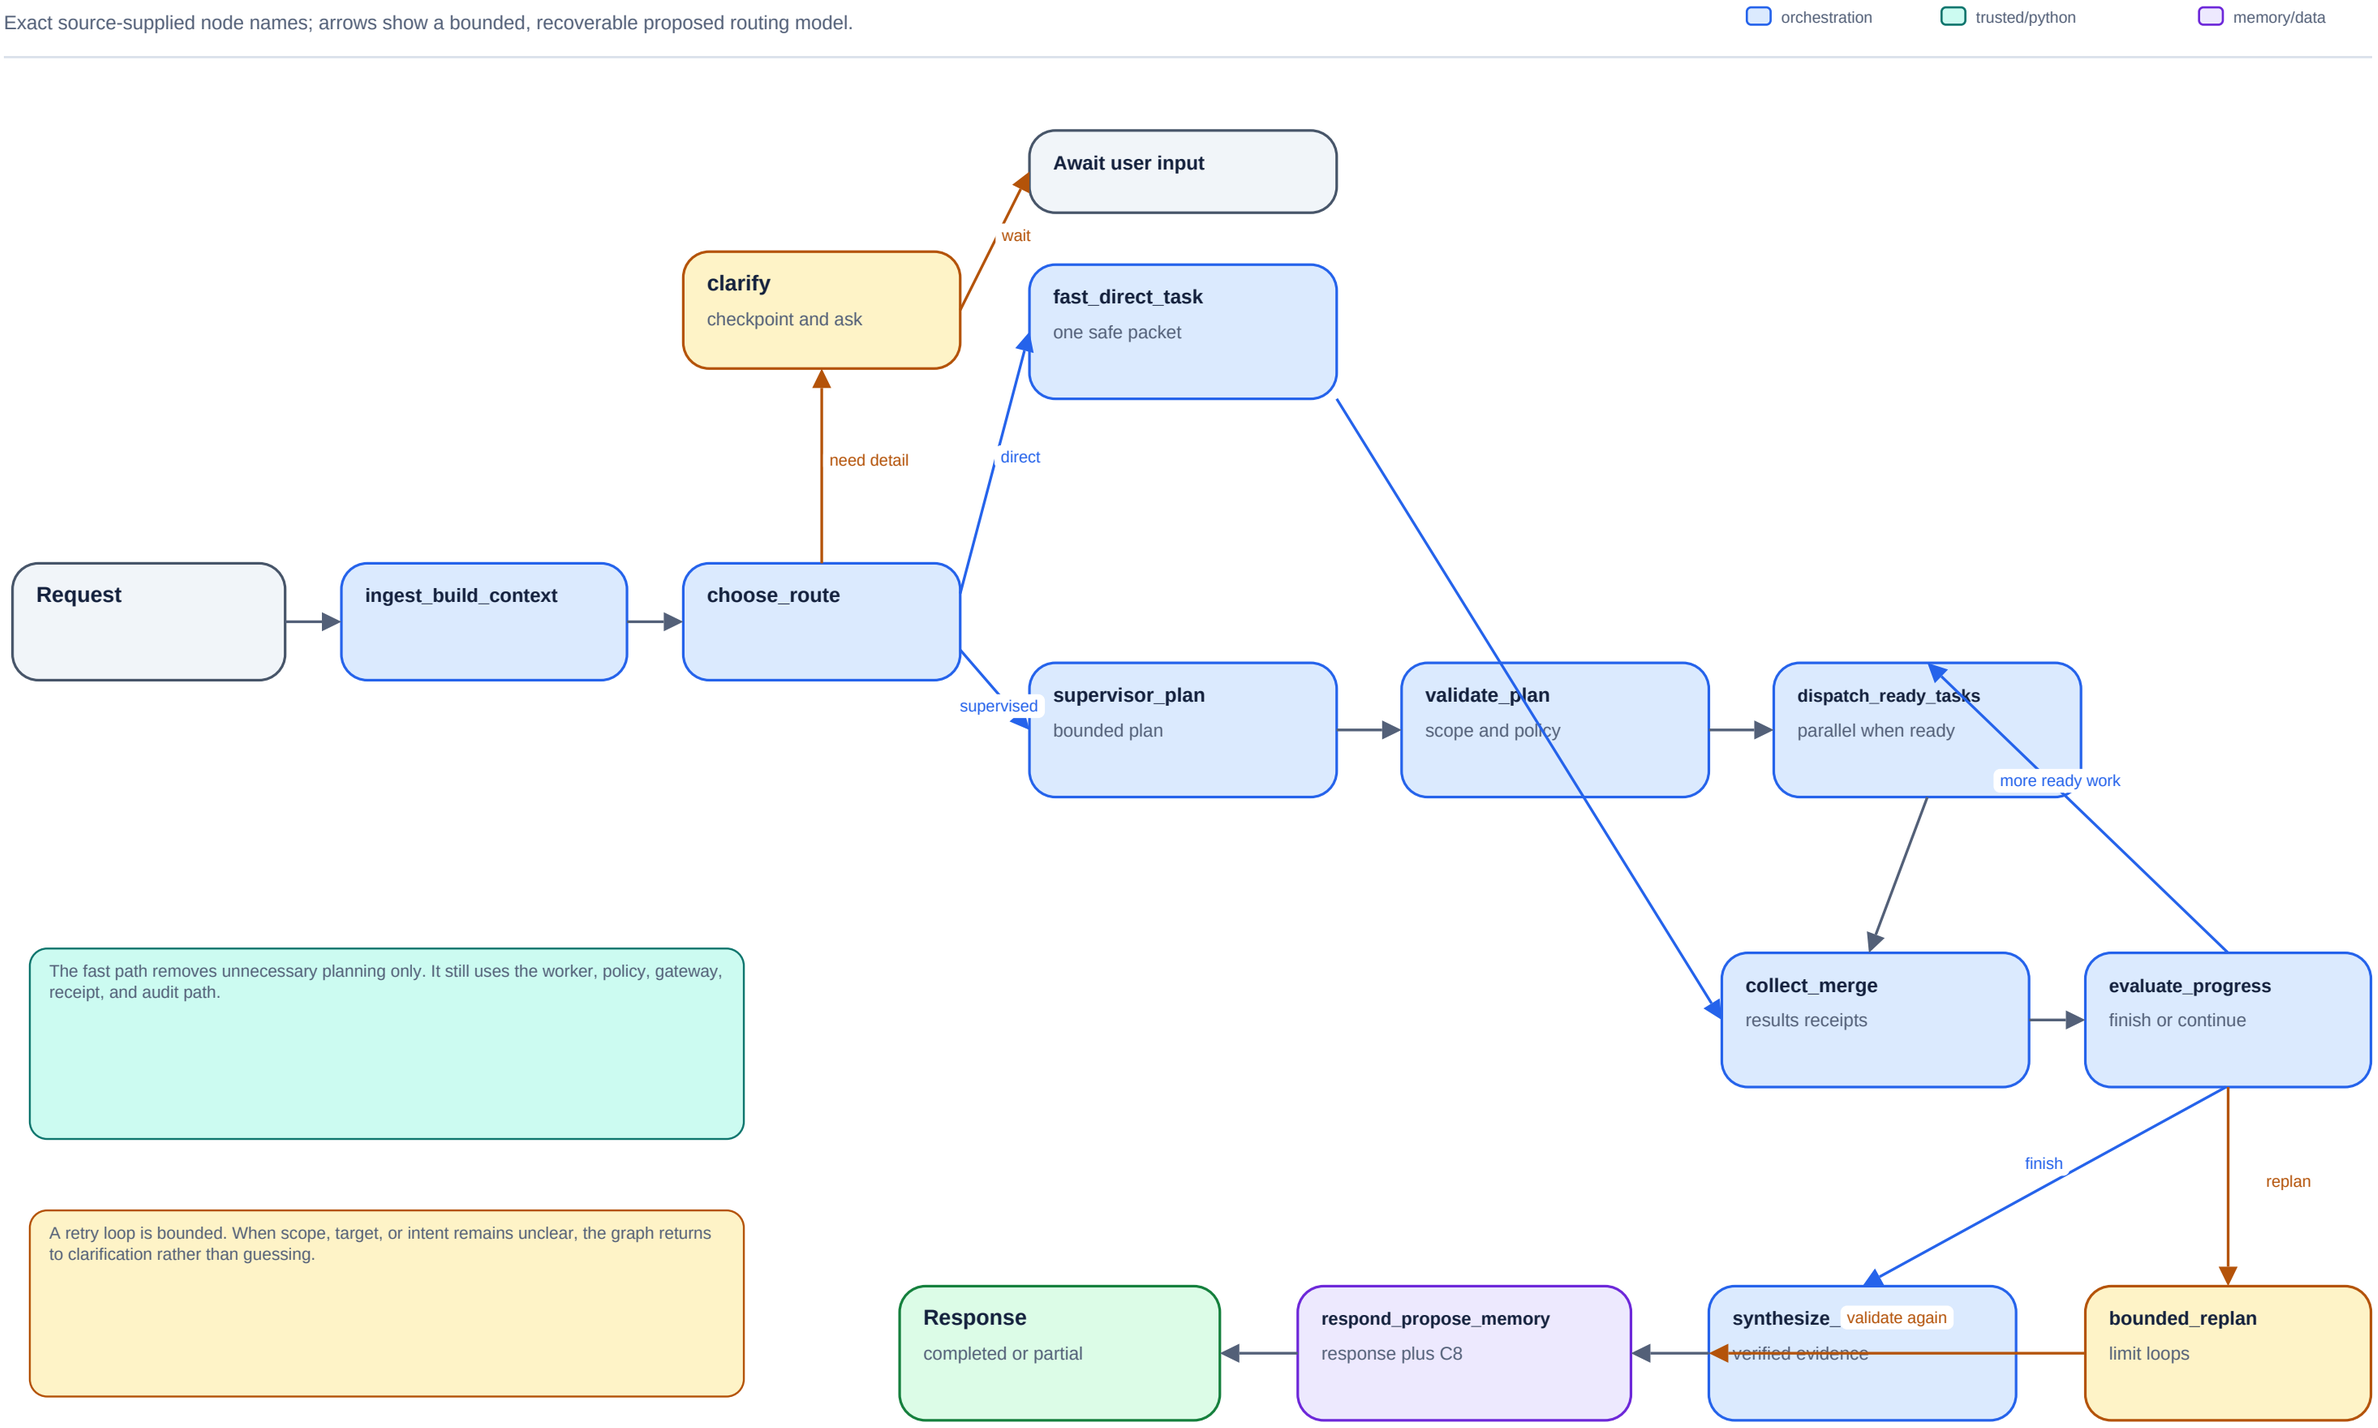

The fast path removes unnecessary planning only. It still uses the worker, policy, gateway, receipt, and audit path.

A retry loop is bounded. When scope, target, or intent remains unclear, the graph returns to clarification rather than guessing.

D3 Agent hierarchy and bounded subgraphs

Specialization improves reasoning, but authority remains in the worker and trusted gateway.

orchestration
trusted/python
memory/data

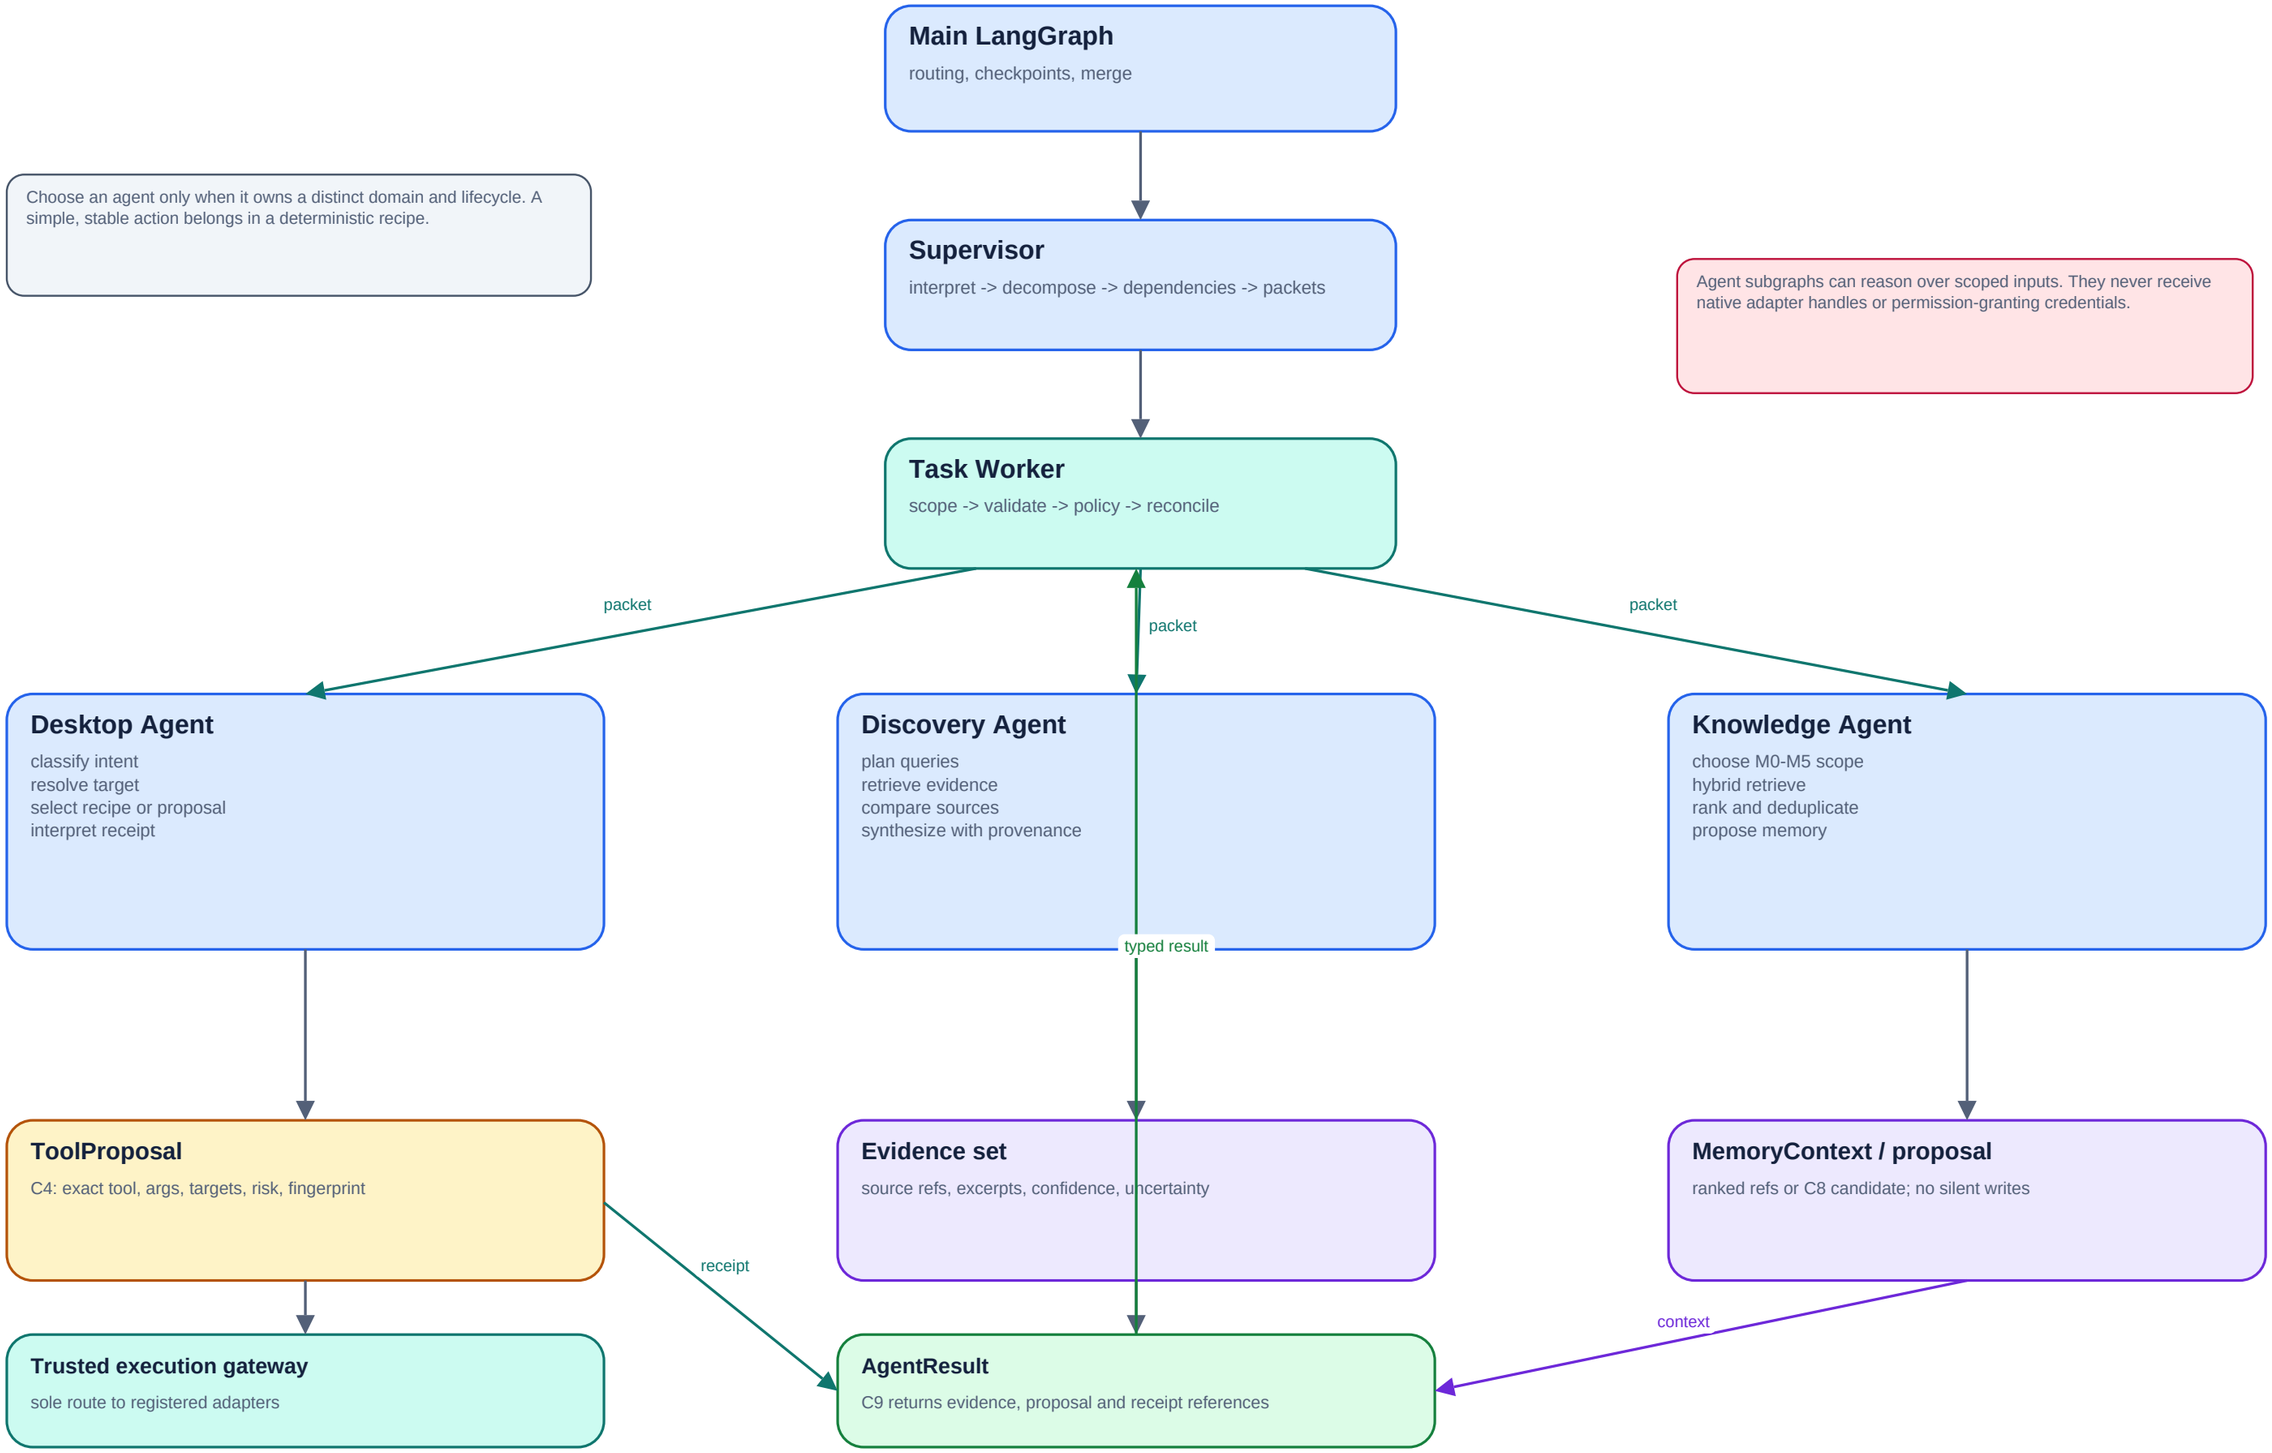

D4 Memory lifecycle and tier selection

M0-M5 supply relevant context and durable knowledge; none of them grants execution authority.

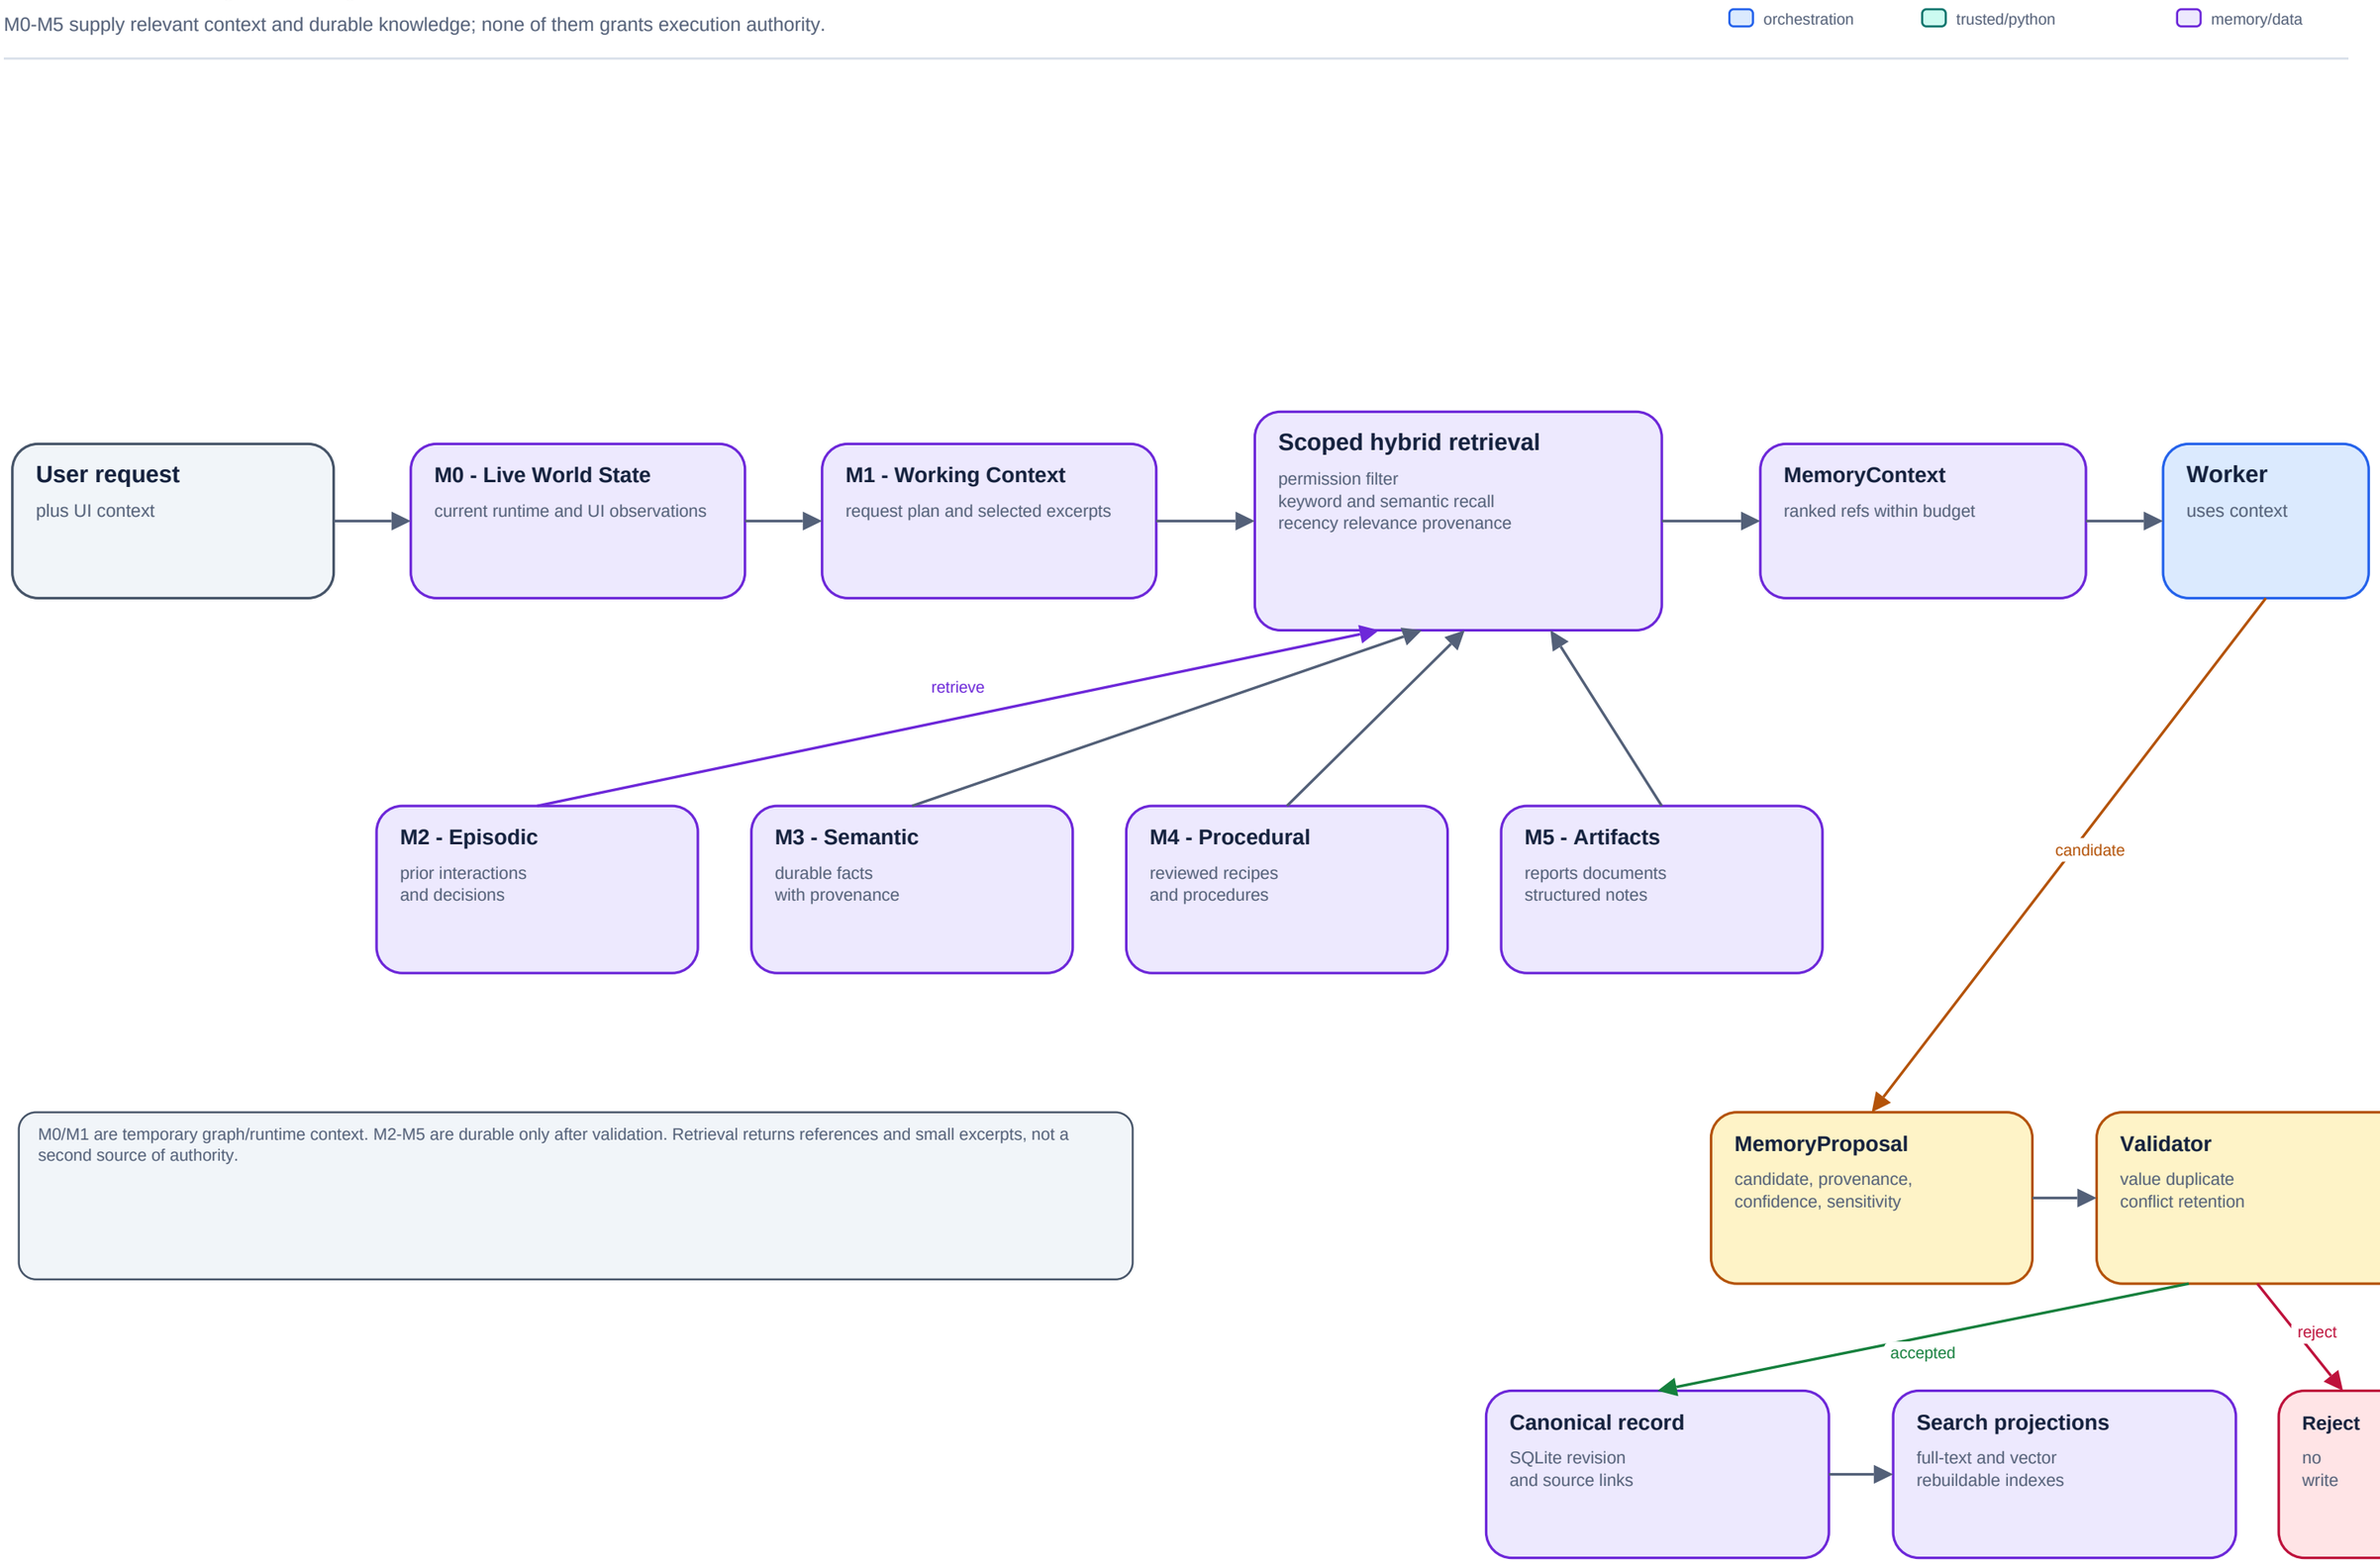

M0/M1 are temporary graph/runtime context. M2-M5 are durable only after validation. Retrieval returns references and small excerpts, not a second source of authority.

D5 Tool approval and trusted execution

A proposal becomes an execution request only after exact validation, policy, and any required human approval.

orchestration | trusted/python | memory/data

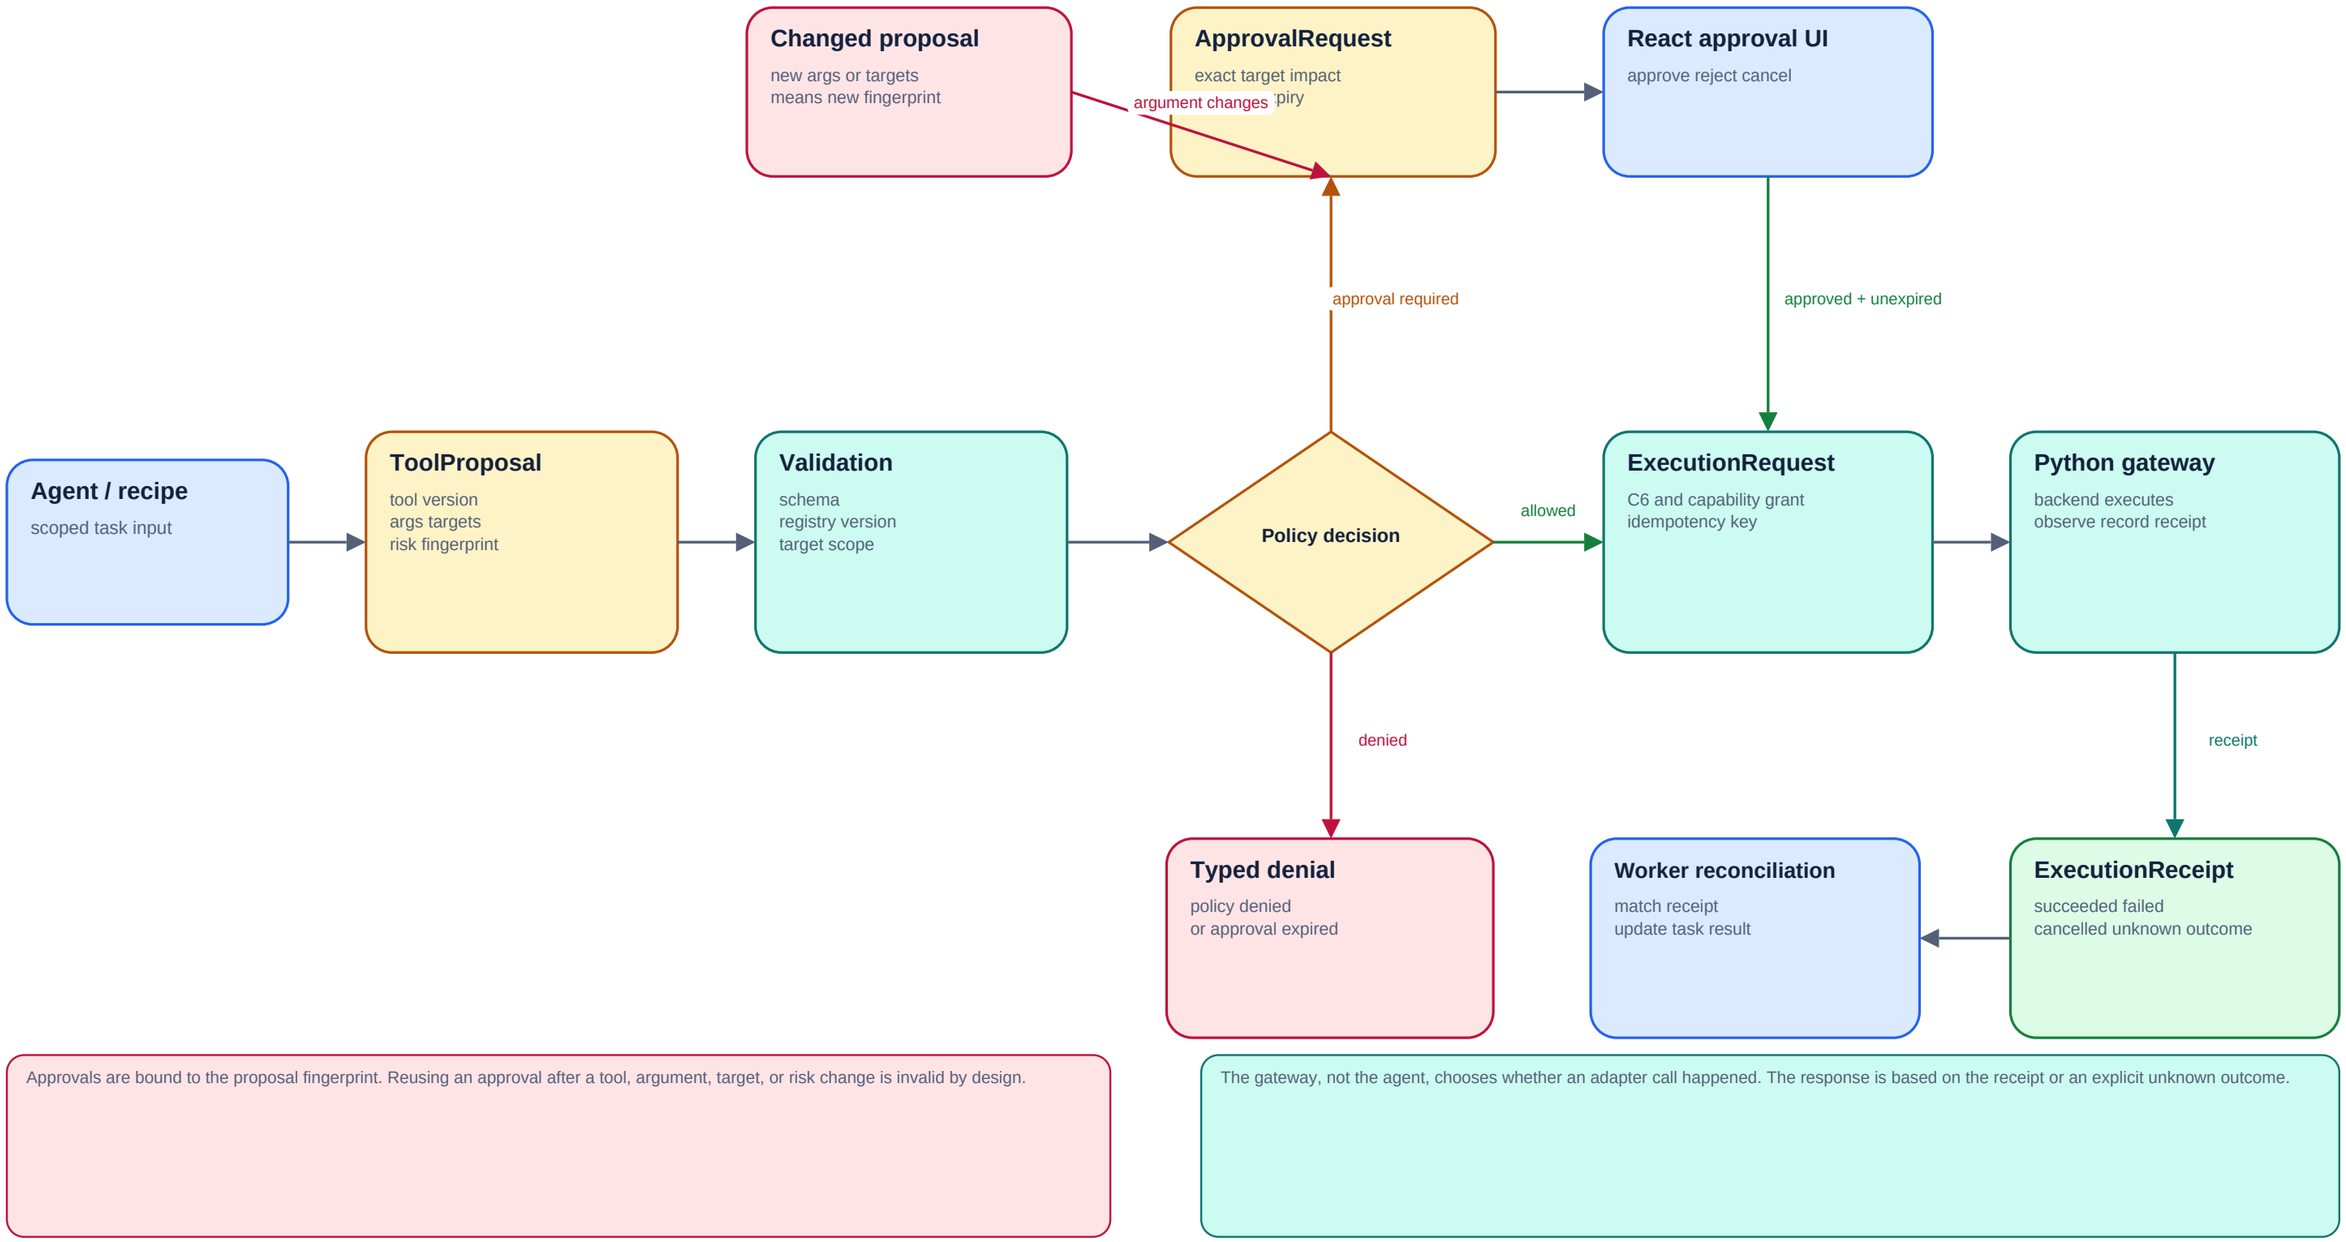

Approvals are bound to the proposal fingerprint. Reusing an approval after a tool, argument, target, or risk change is invalid by design.

The gateway, not the agent, chooses whether an adapter call happened. The response is based on the receipt or an explicit unknown outcome.

D6 Persistence and relationship model

Proposed normalized local-first records. SQLite holds canonical metadata; artifacts and indexes live beside it.
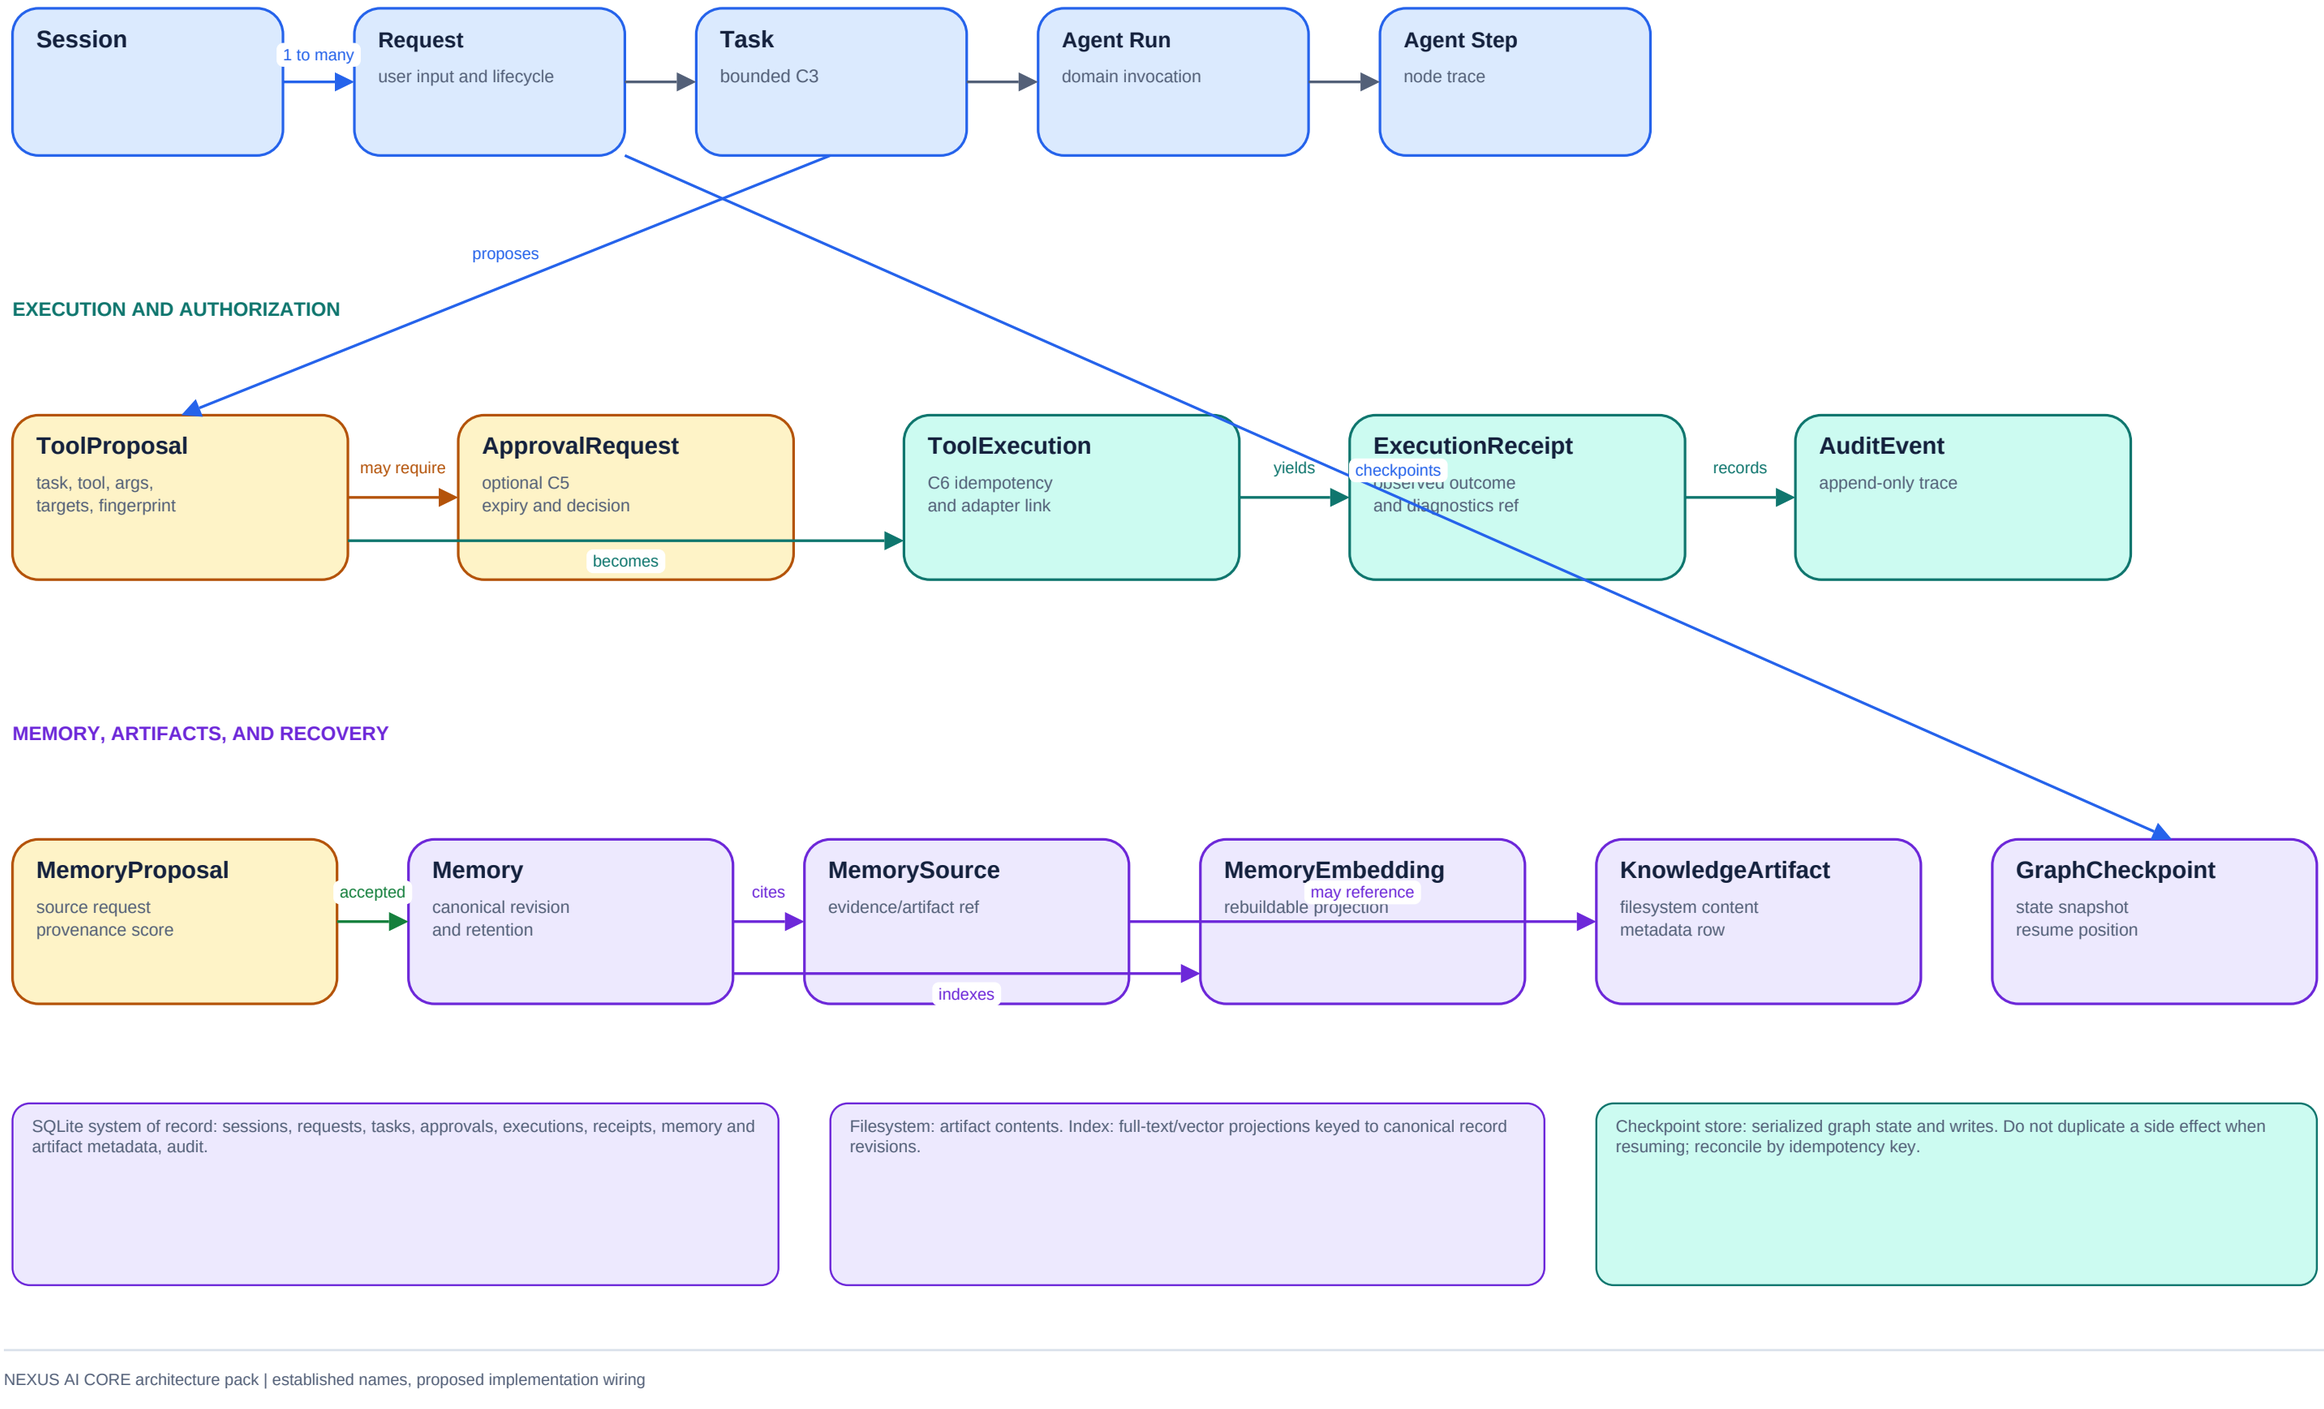
orchestration trusted/python memory/data